




A PROPOSAL FOR PROJECT ARES · RÉMI COUTURE

Eight training games, already playable.

What I have built for Ares on my own initiative, why it teaches, how it fits the platform, and a plan to keep building games for Circadence and its clients.

I have eight working educational games for Ares. I would like to keep making them.

8

games, each playable from briefing to debrief in a browser

209





automated tests, all passing, covering the learning rules of every game

8

different genres: strategy, inbox drill, triage board, graph puzzle, SOC case, dispatch desk, SQL hunt, arcade flight

0

ranges to spin up. Each game is a static bundle that embeds in a course page



What this is

Real training content, not mock-ups: authored scenarios, evidence, policies, queries, and a written explanation for every important mistake. The next pages quote it.

Why it matters to Ares

Missions need a live range. These need nothing. They give learners practice between missions, give sales a demo that runs anywhere, and fill slots the platform has already reserved for games.

What I am asking








Keep me on to take these from prototype to production with the games and content teams, and to keep building new ones. I have set out three ways to do that, including a half-time option.

One skill each. Three to five minutes. A debrief that explains the result.

GAME	PLAYS LIKE	WHAT THE LEARNER PRACTISES	AUTHORED CONTENT
Network Defense Keep Finance Online	Isometric strategy	Routes and dependencies, reading evidence, writing a scoped access rule, backup routes that keep policy	3 rounds, 3 traffic flows, 10 evidence lines, 4 policy lines, 2 reasoning checkpoints, a final no-hint challenge
Phishing Speed Drill The Payroll Request	Timed inbox	Sender, reply-to and real link destination; when to report, mark safe, or verify by phone	8 messages (3 phishing, 3 legitimate, 2 need verification), 32 cues, each with a reason
Incident Triage Race The Busy Morning	Alert board	Prioritising by evidence and business impact, not the severity tag; revising when facts change	6 alerts, 29 timeline evidence items, 18 authored assessments, 2 mid-shift updates
SOC Dashboard Challenge A Strange Login	Investigation case	Telling discriminating evidence from evidence that fits both stories; proportionate response under time	2 hidden truths with identical openings, 9 evidence items in 2 variants, 11 timed events
Cyber Dispatch First Shift	Dispatcher desk	Asking the right questions, giving the safe next step in time, handing off only what is confirmed	6 reports across 2 shifts, 38 questions, 26 facts, 24 guidance options
Query Hunt Unusual Access	SQL investigation	WHERE, GROUP BY, ORDER BY, JOIN on real tables, and not over-claiming from the results	5 challenges, 4 tables, 2 datasets, real SQLite running in the browser
Cloud Command First Orbit	Arcade flight	Capacity through a dependency, cost against availability, least privilege, the IaaS responsibility split	3 checkpoints, 11 evidence items, 12 justifications with feedback, 7 permissions
Privilege Path Puzzle Who Can Read Payroll	Graph puzzle	Tracing nested groups and inherited grants, and making the smallest fix that keeps business access	3 puzzles, 22 nodes, 13 access checks each with a business reason

ALL EIGHT, IN MOTION

Every clip is real play: a decision, a mistake, the explanation, the fix.



The gameplay reel: every game, one mistake, one explanation, one fix.

The screenshot displays the Ares Games interface. At the top left, it shows 'Round 1 / 3 - Establish service' with a goal to 'Restore the connection Finance needs so supplier invoices are delivered.' The top center shows 'Action points 100 +100 / round'. The top right shows service status: Service 100%, Security 100%, and Efficiency 100%. Below this is a 'Needs review' section with three items: 'Needs review' (yellow), 'You judged permitted' (blue), and 'You judged a violation' (red). The main area is a network diagram with nodes: INTERNET, GATEWAY, INSPECTION POINT, CORE SERVER, FINANCE, and OFFLINE. A dashed line shows a path from the GATEWAY through the INSPECTION POINT to the CORE SERVER. A red circle highlights the CORE SERVER. The bottom left has a video player with a play button and a timer '0:00 / 3:15'. The bottom center has buttons for 'Inspect', 'Access rule' (40 points), 'Add route' (60 points), and 'Revise' (10 points). The bottom right has a 'Launch wave' button. On the right side, there is a 'Back to overview' button, an 'ACTION - 60 AP' button, and an 'Add route' button. Below these are two route options: 'A Core → Finance' and 'B Gateway → Finance', each with a 'Build (60 AP)' button. A text box explains that dashed lines show the shortest working path starting from the Gateway.

Round 1 / 3 - Establish service
Restore the connection Finance needs so supplier invoices are delivered.

Action points
100 +100 / round

Service 100%
Security 100%
Efficiency 100%

Needs review You judged permitted You judged a violation

Planning — simulation paused

INTERNET GATEWAY INSPECTION POINT CORE SERVER FINANCE OFFLINE

EXTERNAL DMZ - INSPECTED INTERNAL

Inspect 0:00 / 3:15

Access rule 40 Unlocks in round 2 / 2

Add route 60 Connects services 3

Revise 10 Change a decision 2

Launch wave

Back to overview

ACTION - 60 AP

Add route

Dashed lines on the board show where a connection can be built. Traffic takes the shortest working path, and every path still starts at the Gateway.

A Core → Finance
Connects Finance to the Core server inside the internal zone.
Build (60 AP)

B Gateway → Finance
Connects Finance straight to the Gateway — the shortest path from the internet.
Build (60 AP)

What to look for

- The learner makes a real decision with real content.
- The game explains a mistake in one or two specific sentences.
- The learner fixes it and sees the result change.


Also in this folder

Ares-Games-Gameplay-Reel.mp4 to share on its own, and the **play** folder with all eight games. Every game slide has a Play button.

Eight interfaces, one visual language. Dense when it needs to be, never confusing.




Rule editor with a live preview of who the rule allows and denies




Before and after table that proves a fix kept business access




Flight between stations, collecting evidence on the way



“What changed” coach note, with rewind to any decision



Decision replay rebuilt from the learner's own moves





Habits, not just a score: false alarms, misses, skipped checks



Four-part score with response time kept separate



Feedback built from the simulation's own numbers



How it is scored. 40% service delivered, 40% prohibited activity denied, 20% efficiency. Two reasoning checkpoints and the final challenge are recorded separately as learning evidence.

Proved by tests. “Blanket blocking cannot earn the best result” and “backup routes preserve the access policy”, each checked on 8 seeds. 32 tests.

Depth and replay. A new seed changes the traffic order, the attacker’s address and service, and which link fails.

GAME 1 OF 8 · ISOMETRIC STRATEGY

Network Defense: Keep Finance Online

Teaches: How routes and dependencies keep a service up, how to read evidence, and how to write an access rule that stops the attacker without stopping the business.

EVIDENCE ON SCREEN

“Source 198.51.100.23 is not on any approved list and is outside the company network.” “Destination: Finance — Remote desktop admin (RDP 3389).” “Pattern: repeated sign-in attempts, each with a different username.”

AFTER A COMMON MISTAKE



“You connected Finance straight to the Gateway. Invoices reached Finance, but Finance couldn’t reach its database on the Core server, so each one failed. Services depend on other services — a path from the internet isn’t enough.”

FINAL CHALLENGE, NO HINTS

“Which single rule best fits the policy?” Answer: DENY 203.0.113.200 → HR : Admin console

Why it sticks. The learner watches a rule succeed or fail on live traffic, so cause and effect are remembered together.

Where it fits Ares. Fits the reserved packet-gate and network-segmentation slots. Firewall and ACL rules are listed as untapped.



How it is scored. 100 for the best response, +20 per strong cue cited, a streak bonus, and a speed bonus that only counts after a correct decision. A false alarm scores zero.

Proved by tests. Always-report, always-safe and always-verify each “stay out of the top two ratings”. “Slow but correct play still earns the top rating.” 53 tests.

Depth and replay. 8 messages, 32 cues. The first message is guided and untimed; later batches are shuffled per seed.

GAME 2 OF 8 · TIMED INBOX DRILL

Phishing Speed Drill: The Payroll Request

Teaches: Reading the sender, the reply-to and the real link destination, backing a call with evidence, and knowing when only a phone check can settle it.

A LOOK-ALIKE DOMAIN

helpdesk@acme.example.it-notices.example “A domain is read from the right. This address belongs to it-notices.example; ‘acme.example’ is only a label in front of it.”

A GENUINE ADDRESS, STILL NOT SAFE



A bank-change request from a real colleague’s address: “The address is genuine, but a matching address doesn’t prove who is typing. Mailboxes get compromised.”

AFTER A FALSE ALARM

“This was a real payslip notice. External mail and a generic greeting are not evidence by themselves; the domain, link, and timing all matched. False alarms slow down real work.”

Why it sticks. Citing the evidence before deciding builds the habit of looking, not guessing.

Where it fits Ares. Fits the reserved phish-sprint slot. Suits awareness training for every employee, not only analysts.



How it is scored. Points for the best or a defensible lane, a matching reason, revising on new evidence and correct links. Missed criticals and needless escalations cost points. Time is reported separately.

Proved by tests. “Never treats the severity tag as a valid justification.” “Rewards revising when evidence changes, not standing still.” 13 tests.

Depth and replay. 6 alerts, 29 evidence items, 2 updates over 3 rounds, with a 7-moment replay in the debrief.

GAME 3 OF 8 · ALERT TRIAGE BOARD

Incident Triage Race: The Busy Morning

Teaches: Setting priority from evidence and business impact, linking alerts that are one incident, and revising when new facts arrive.

TAGGED “LOW” IN THE QUEUE

A new device on an account that “approves outgoing payments”. “The quietest alert in the queue turns out to be an attacker with a working sign-in on the payments portal. Business impact, not the severity tag, sets the urgency.”

A MID-SHIFT UPDATE

“Infrastructure confirms change ticket CHG-4471: backups moved to a new offsite storage provider this morning... Nothing else in the queue is explained by it.”

A JUSTIFICATION THAT ALWAYS SCORES ZERO

“The severity tag says so.”

Why it sticks. Being wrong about the quiet alert, then seeing why, is the lesson people repeat to colleagues.

Where it fits Ares. Matches the design guide’s “Alert Triage” concept. A natural team game for workshops.

PROJECT ARES | SOC TRAINING PRACTICE

DECISION 1 / 3

UNTIMED

ACTIVE CASE

Suspicious Account Activity

Unusual activity detected for a finance user with possible data exfiltration.

User

t.nguyen

Asset

FIN-LAP-028

Location

New York, US

First seen

Today 10:14

Case clock

10:33

CASE STATUS

Investigating

Login activity

PARTIAL SIGNAL

Cite

Multiple logins from unusual locations.

10:14 Successful login New Yor...

10:27 Successful login London, ...

10:31 MFA challenged London, ...

10:32 Successful login London, ...

+ Device and session detail

2 min

+ MFA record

2 min

+ Travel record and manager check

4 min

Endpoint event

PARTIAL SIGNAL

Cite

Process and file activity on FIN-LAP-028.

10:22 Uncommon process powershell.exe

10:24 Archive created finance_q3.zip

10:25 Large file access 1.2 GB (243 files)

10:26 Process network connection to 203.0.113.42

+ Process tree

2 min

+ File access detail

2 min

+ Change ticket lookup

2 min

Outbound traffic

PARTIAL SIGNAL

Cite

Data transfer to an uncommon external IP.

Top destination

203.0.113.42

Uncommon destination

Total transferred

0.56 GB

Time window

10:26 - 10:33

+ Destination ownership and history

2 min

+ Live transfer status

2 min

+ Other hosts using this destination

2 min

Evidence timeline

Key events in this case (chronological)

10:14 Login New York

10:24 Archive created finance_q3.zip

10:26 Outbound connection 203.0.113.42

10:31 MFA challenged London

10:32 Login successful London

10:33 Now Case clock

Your decision

Choose the best response for this situation.

Evidence requests 1

Your reasoning

Select the primary reason, and cite up to 3 pieces of evidence above.

Debrief preview

Your decision + your reasoning above.

100

100%

This was a compromise.

Lead analyst

Business consequences

Decision replay

How it is scored. Outcome 40, decisions 36, reasoning 24. “Each call is judged on what you knew at the time, not on how the case turned out.”

Proved by tests. “No single blind opening wins in both variants.”

“Investigating first is not always right: delay lets the incident progress.” 13 tests.

Depth and replay. The seed decides whether it is a compromise or approved work. Practice mode can rewind to any decision.

GAME 4 OF 8 · SOC INVESTIGATION CASE

SOC Dashboard Challenge: A Strange Login

Teaches: Separating evidence that discriminates between explanations from evidence that fits both, and responding in proportion while data is leaving.

SAME OPENING, TWO POSSIBLE TRUTHS

Compromise: “OUTLOOK.EXE → WINWORD.EXE (Invoice_Q3_revised.docm) → powershell.exe -WindowStyle Hidden -EncodedCommand ...” Approved work: “Task Scheduler → powershell.exe - File QuarterClose-Export.ps1”.

EVIDENCE THAT PROVES NOTHING


“A passed prompt shows that someone tapped Approve. Intruders get prompts approved too, so this fits both explanations.”

AFTER OVER-REACTING

“You had strong evidence that this was approved work and contained anyway. The business paid for an interruption the evidence did not support.”

Why it sticks. Two truths with the same opening teach that the right call comes from evidence, not from matching a familiar story.

Where it fits Ares. A step up from triage toward analyst missions, and a warm-up before a range investigation.



How it is scored. Four parts of 25 per report: evidence gathered, priority, guidance and handoff. Response time is shown but kept out of the score.

Proved by tests. “Escalating everything as high scores poorly.” “Asking every question before acting is not the best run.” 24 tests.

Depth and replay. 6 reports over 2 shifts, 38 questions, 24 guidance options. The second shift changes which report is benign.

GAME 5 OF 8 · DISPATCHER DESK

Cyber Dispatch: First Shift

Teaches: Asking the questions that matter, giving the safe next step before harm lands, and handing off only what the evidence confirms.

A REPORT COMES IN

“Hi! I just got an MFA prompt but I’m not signing in right now. I’m at my desk. Is this legit?” Asking reveals: “The prompt shows a Mail sign-in from Rotterdam. She is in the Lisbon office.”

THE SAFE STEP, IN TIME

“Turn off Wi-Fi and unplug the network cable now. Leave the laptop switched on.”

A HARMFUL STEP AND ITS CONSEQUENCE

“Reply to the email and ask them to confirm the new details.” → “The sender ‘confirmed’ their own request, and Elena queued the payment.”

Why it sticks. Recoverable harm: the learner sees the damage a bad instruction causes, then gets to put it right.

Where it fits Ares. Covers first response and handoff, part of the incident response lifecycle listed as untapped. Good for help desks.

Evidence 3
Not collected yet

Evidence 4
Not collected yet

Evidence 5
Not collected yet

Practice sandbox

Schema

Review material

Reduced motion: off



Gameplay concept • Scenario content pending SME review

CHALLENGE 1 OF 5 • AUTHENTICATION LOGS

An alert fired for the period 3 April 2024, 00:00 to 06:00. List the failed logins in that period.

Result must include event_id, ts
Order Required: ts ascending (oldest first).

Guided blocks

Write SQL

Hint (3 left)

✓ YOUR QUERY, CHECKED AGAINST THE EXPECTED RESULT

```
SELECT event_id, ts, account_id
FROM logins
WHERE result = 'failure'
AND ts >= '2024-04-03 00:00:00' AND ts < '2024-04-03 06:00:00'
ORDER BY ts
```

Runs in your browser on synthetic data. Read-only.

Run query

Event ID	Account ID	Account Name
L-0014	A-1001	John Doe
L-0016	A-1001	John Doe
L-0017	A-1001	John Doe
L-0018	A-1001	John Doe
L-0019	A-1001	John Doe
L-0020	A-1001	John Doe
L-0021	A-1001	John Doe
L-0022	A-1001	John Doe
L-0023	A-1001	John Doe

Which of these does your result show?

Failed logins in the alert window

logins table

Period3 April 2024, 00:00 to 06:00

Failed attempts19

Accounts involved4

19 failed logins fall inside the alert period, across 4 accounts.

maria.gomez: 4 failed logins at 08:55

logins table

AccountA-1001

Time2024-04-03 08:55:00

Several failures in a row on one account.

An account was taken over

no source

BasisFailed logins

Someone broke into an account during the alert period.

Pin evidence

Query Hunt: Unusual Access

Teaches: Filtering, aggregating and joining real log tables, reading the result correctly, and telling a suspicious observation from a confirmed conclusion.

A QUERY THE LEARNER PRODUCES

SELECT account_id, COUNT(*) AS attempts FROM logins WHERE result = 'failure' ... GROUP BY account_id ORDER BY attempts DESC

WHEN THE LEARNER JUMPS AHEAD

“Repeated failed logins alone do not establish compromise: they show attempts. Which evidence shows what happened after them?”

THE DEBRIEF NAMES THE LIMITS

“The records do not show who was at the keyboard, how the password was obtained, or that data left the device.”

Case closed: Unusual Access

1000 / 1000

4000

2500

2500

1500

0

WHAT THE EVIDENCE SUPPORTS

The records show 12 failed logins on alert area from 2024-03-03 00:00:00 to 2024-03-03 05:59:59. One successful login from the game server at 2024-03-03 08:55:00. The user registered, then 8 attempts to reconnect that failed. The user then logged in successfully with a confirmed company.

3. Failed logins in the alert window

Period: 3 April 2024, 00:00 to 06:00

Failed attempts: 19

Accounts involved: 4

19 failed logins fall inside the alert period, across 4 accounts.

4. Maria Gomez: 4 failed logins at 08:55

Account: A-1001

Time: 2024-04-03 08:55:00

Several failures in a row on one account.

5. An account was taken over

Basis: Failed logins

Someone broke into an account during the alert period.

YOUR QUESTIONS

1. Find the failed logins

2. Count attempts by account

3. What happened after the alert?

MISSED INFORMATION

No unreported data on the way. Entry being blocked from a server.

Who was using the device. The log record on the server was deleted.

Who the correct password was entered.

Whether the compromised log left the device.

Why also data was deleted from the server at 2024-03-03 08:55:00. And the account name.


How it is scored. Out of 1,000: correct results, evidence cards, interpretation and conclusion. A right conclusion without its evidence scores zero. Hints are counted, never deducted.

Proved by tests. “Accepts equivalent queries written differently.” “Never grants evidence for running a query.” Real SQLite runs in the tests. 31 tests.

Depth and replay. A second dataset makes the highest failure count the wrong account to follow. Guided tiles or written SQL.

Why it sticks. Writing the query is retrieval practice. Being stopped from over-claiming teaches analytic discipline.

Where it fits Ares. Fits the reserved log-hunt slot. Log analysis is listed as untapped.



How it is scored. Continuity 35%, policy 30%, resources 15%, justification 20%. A failed access review marks the security objective failed whatever the score.

Proved by tests. “Adding capacity helps only where capacity is the constraint.” “Does not let a high score hide a failed security objective.” 18 tests.

Depth and replay. 3 checkpoints, 11 evidence items. Fly manually or on autopilot; flying is reported apart from the learning score.

GAME 7 OF 8 · ARCADE FLIGHT

Cloud Command: First Orbit

Teaches: Finding the real capacity constraint through a dependency, balancing availability against cost, least privilege, and who is responsible for what on IaaS.

REQUIREMENTS

“Serve at least 99% of checkout requests at the peak.” “Keep 10% headroom above the expected peak on every service in the request path.” “Cost ceiling: 120 credits for the 2-hour surge window.”

AFTER SCALING THE WRONG SERVICE

Checkout “has enough capacity, but each request calls” the Payments API, “which can serve only” 1,200 rps. More checkout instances “cannot help until” it has capacity.

LEAST PRIVILEGE

A backup account holds 7 permissions. “Reading a snapshot requires keys:use on that key. The ticket does not mention it.”

Why it sticks. The numbers in the feedback come from the simulation, so the explanation matches exactly what the learner did.

Where it fits Ares. Cloud security is listed as untapped. The arcade layer shows the same content can carry a lighter style.

GAME 8 OF 8 · ACCESS GRAPH PUZZLE

Privilege Path Puzzle: Who Can Read Payroll

Teaches: Tracing nested groups and inherited grants to find every route to a file, then making the smallest change that closes it without breaking business access.

THE BRIEF

"An engineer opened a payroll file. Nobody remembers giving them access." Hidden in the graph: "Added during a payroll deadline two years ago so Finance could get in quickly."

A CHECK THE FIX MUST STILL PASS

"svc-backup · Write · Ledger — must be denied": "A backup job copies data. It never needs to change it."

AFTER A WRONG EXPLANATION

"Signing in is authentication: it proves who Alex is. What Alex may open is authorization, and that came from a grant."

Why it sticks. Predict, test, then explain: the learner commits to an answer before seeing the result.

Where it fits Ares. Matches the design guide's "IAM Auditor" concept. Identity and access is listed as untapped.

PROJECT ARES / TRAINING MINIGAME

PRIVILEGE PATH PUZZLE

Scenario 1 / 3: Who Can Read Payroll?

An engineer opened a payroll file. Nobody remembers giving them access.

OBJECTIVE

Stop Alex reading Payroll Files without breaking Finance's access.

CHANGES

SCORE

Member of

Grant

Contains

Not a permission

Permission Inspector

[Access rules](#)

FROM

All Employees

Group

TO

Payroll Files

\\files\payroll

GRANT

Read

Members of All Employees may read Payroll Files and anything inside it.

Added during a payroll deadline two years ago so Finance could get in quickly.

NARROW THIS GRANT

Granted to **All Employees**

☒ Read

Remove grant

Access checks Select one, predict the result, then test it.

Alex Chen · Read · Payroll Files

Dana Ortiz · Read · Payroll Files

Alex Chen · Read · Company Handbook

Test access

Trace the selected

Verify fix

[illegible]

How it is scored. 100 for solving, a bonus for matching the minimum number of changes, points for correct predictions, and 40 for explaining it right first time.

Proved by tests. “Never treats connectivity as authorization.” Each puzzle “cannot be solved with fewer changes than par” and “fails when access is removed for everyone”. 25 tests.

Depth and replay. 3 puzzles, 22 nodes, 13 access checks. Inspect, predict, test, change, verify, explain.

Each game is built on how people learn, not on points and badges.

RETRIEVAL PRACTICE

The learner produces the answer

They write the SQL, build the rule, place the alert, choose the fix. Producing an answer builds memory that reading and recognising do not.

IMMEDIATE, SPECIFIC FEEDBACK

Every mistake explains itself

Feedback names the decision and the reason, at the moment it happens. Network Defense alone has about 19 explanation templates.

SAFE FAILURE

Consequences without damage

A wrong call queues a fraudulent payment or cuts off a working laptop, in fiction. The learner sees the cost, then gets to fix it.

JUDGMENT, NOT RECALL OF FACTS

Evidence before conclusion

Inspecting never classifies for you. Weak evidence never counts as proof. A correct conclusion without its evidence scores zero in Query Hunt.

REFLECTION

A debrief that replays your decisions

Choice, observed result, why, and how it connects to the mission. Network Defense rebuilds the exact moment from the command log.

SHORT AND REPEATABLE

Five minutes, different next time

Seeds and variants change traffic, attackers, the failing link, or which report is benign, so a replay is practice and not memorisation.

One rule runs through all eight: **the game score and the learning evidence are recorded separately**. As one debrief puts it, “this record is evidence, not a grade.” That is what makes the result usable for reporting.

The hard part of a learning game is balance. I test it the way a game studio would.

If a learner can win by doing the same thing every time, the game teaches that habit. So every game has tests that prove the lazy strategy loses.

LAZY STRATEGY	WHAT THE TESTS PROVE
Block everything	"Blanket blocking cannot earn the best result", checked on 8 seeds
Report every email	Always-report, always-safe and always-verify each "stay out of the top two ratings"
Escalate everything	"Penalizes escalating everything alarming" and "penalizes monitoring everything"
Trust the label	"Never treats the severity tag as a valid justification"
Always investigate first	"Investigating first is not always right: delay lets the incident progress"
Go fast	"Slow but correct play still earns the top rating"; speed only counts after a correct decision

The first ten seconds

A guided, untimed first step in each game. The learner acts before they read a manual, the way a good game tutorial works.

An economy that forces trade-offs

Action points, analyst slots, credits and case minutes are all limited. Limited resources are what turn knowledge into decisions.

Feedback you can feel

Routes light up, tokens bounce off a rule, a service goes dark. Shape and icon back up colour, so the signal never depends on colour alone.

A reason to play again

Ratings, a clean-run target, and a variant that changes the right answer. Query Hunt's second dataset makes the highest count the wrong account to follow.

These are engineered to be reviewed, tested, embedded and extended.

Content lives in one file per game

Scenarios, evidence, policy and explanations sit in a single data file, apart from the code. A content developer or SME can review and change the teaching without touching the engine.

The rules run without a browser

Each game's simulation is pure logic, separate from the screen. That is why 209 tests can check the learning rules in under two seconds.

Deterministic and replayable

The same seed and the same commands always give the same result. It powers the debrief replays, and it would let the platform verify a result on the server.

A host contract already written

Each game posts started, statement, completed and navigate messages to its host page, with an activity ID per scenario. A draft integration document describes the payloads.

SME review packs ready to sign

Four games come with a review sheet that lists every objective, simplification and explanation for an expert to approve, correct or reject.

Accessible by design

Keyboard play, focus-trapped dialogs, screen-reader announcements, and reduced-motion and larger-text settings. Layouts stack for small screens.

~22,000

lines of TypeScript, including tests

8 pages

one static page per game, one build

0 servers

no backend, no range, no sign-in

AI-assisted


I design and direct, and build with AI tooling. That is how one person ships this much.

With these, the Games Hub goes from 6 games to 14.


New · my eight, playable now. Judgment under pressure: what to do with incomplete evidence, limited time and a business to keep running.




Network Defense
Routes, evidence, scoped access rules




Phishing Speed Drill
Sender, link and reply-to evidence




Incident Triage Race
Priority from evidence and impact



SOC Dashboard Challenge
Discriminating evidence, proportionate response



Cyber Dispatch
First response and clean handoff



Query Hunt
SQL log investigation



Cloud Command
Capacity, cost, least privilege



Privilege Path Puzzle
Groups, grants and minimal fixes

In the hub today · built by the games team. Framework fluency: tactics, the kill chain, patterns, ports and attack paths.



Astro-Defense
ATT&CK tactics and kill chain responses



Cylitaire: Kill Chain
Kill chain order and tools



Trivialoot
Six-category knowledge board



RegExile
Regular expressions



Portflow
Ports and protocols



Threat Modeler
Attack paths and detection coverage

These add to the Games Hub. They cover ground the platform has already marked as open.

The six games in the hub today, built by the games team, cover ATT&CK tactics, the kill chain, regex, ports, trivia and attack paths. The team’s own design guide lists what is still uncovered, and the router already holds five reserved game slots with nothing behind them.

OPEN IN THE DESIGN GUIDE OR ROUTER	WHAT I HAVE PLAYABLE
Reserved slot phish-sprint	Phishing Speed Drill
Reserved slot log-hunt · “Log Analysis / SIEM: entirely untapped”	Query Hunt
Reserved slots packet-gate, network-segmentation-sim · “Firewall / ACL Rules: entirely untapped”	Network Defense
Tier 2 concept “Alert Triage”	Incident Triage Race, SOC Dashboard Challenge
“Cloud Security (IAM, exposure): entirely untapped”	Cloud Command, Privilege Path Puzzle
“Incident Response Lifecycle: entirely untapped”	Cyber Dispatch (first response and handoff)

Different from what exists

The hub’s games drill recall of frameworks. These drill judgment: what to do with incomplete evidence, limited time and a business to keep running. The two sets complement each other.

They line up with the team’s approval gate

The mechanic collapses without domain knowledge. Accuracy outranks speed. Feedback happens during play. Replay is driven by content. Explicit NICE work-role and KSA mapping is the piece still to add, with the content team.

Small integration step

The hub uses a ready, score-update, complete message flow. My games already post the equivalent events, so each needs a thin adapter, not a rewrite.

Eight genres so far. The method is what scales, not one look.

How a new game gets made

- **Start from the skill.** One learning objective, one decision the learner keeps getting wrong.
- **Find the mechanic that is the skill.** If you can win without the knowledge, it is the wrong mechanic.
- **Write the content as data.** Scenario, evidence, and an explanation for each mistake, in one reviewable file.
- **Prove the balance.** Tests that show the lazy strategies lose and careful play wins.
- **Hand the SME a review sheet.** Every objective and simplification, ready to approve or correct.

The shared pieces already exist: the seeded random generator, the replay log, the host messages, the accessibility helpers and the build. A new game is a data file, a simulation, a screen and a test file.

What I would build next

CONCEPT	WHY
IR Sequencer	Tier 1 in the design guide. Order the response steps under pressure.
CSF Incident Sorter	Tier 1. Route events to Identify, Protect, Detect, Respond, Recover.
ATT&CK Technique Mapper	Tier 1. Takes the hub from tactic level to technique level.
More cases for the eight	New scenarios are content, not code. Query Hunt Case 02, a second phishing inbox, a harder shift.
Team modes for workshops	Head-to-head triage and dispatch for Security Immersion Workshop events, with a shared leaderboard.
Mission warm-ups	A two-minute game before a range mission that drills the one concept the mission depends on.

Games are the cheapest way to add practice, reach and proof to the platform.

Practice that costs nothing to run

A mission allocates live virtual machines. A game is a static page. Learners can practise as often as they like without a range starting up.

A demo that runs anywhere

No sign-in and no install. A prospect can play a round in a meeting, on a booth screen, or from a link in an email, and understand what Ares teaches.

More for every licence

New games widen the catalog between major releases, and give customer admins something fresh to assign.

Content for workshops and events

Short, scored, spectator-friendly games suit Security Immersion Workshops, where dozens of people play at once from a browser.

Evidence for reporting

Each game records decisions and reasoning separately from the score. Once mapped to work roles, that feeds readiness reporting with more than a completion tick.

Ready for the Japanese market

Text is held as data, so games can be localised the same way as the platform. I have already worked on the Japanese strings across the player app.

I have not measured learner outcomes yet, and I am not claiming numbers. The first thing I would do with the team is put these in front of real learners and measure completion, replays and what they get wrong.

They are prototypes. Here is exactly what stands between them and production.

OPEN ITEM	WHAT IT NEEDS
SME review	All content is marked draft. Four games have review sheets ready; four still need one written. I need a few hours of expert time per game.
Learner playtesting	None yet. Internal staff first, then a friendly customer cohort.
Framework mapping	The games do not yet name NICE work roles or KSAs. That mapping should be done with the content team so it matches the catalog.
Platform integration	The host contract is a draft that has not been agreed with the platform team. Course and mission text is mocked.
Localisation	English only today. Japanese follows SME sign-off.
Borrowed art and patterns	Network Defense uses sprites reused from another developer's project with their go-ahead. Before shipping I would confirm that in writing or replace the art.
Depth	Three games have seeded variation, two have a second variant, three have one scenario. More cases are content work.

DEBRIEF - CASE 01 - DATASET A

Case closed: Unusual Access

Investigation score
★ 1000 / 1000

Correct results
400 / 400

Evidence selection
250 / 250

Interpretation
200 / 200

Conclusion
150 / 150

Hints used
0
Reported separately. No points deducted.

WHAT THE EVIDENCE SUPPORTS

The records show 12 failed logins on alex.chen from 2024-04-03 (Singapore), then a successful login from the same address at 02:21 on a device that is not registered, then 6 accesses to restricted files in /finance/payroll. Together this is consistent with unauthorized use of the account and justifies escalation. It is not confirmed compromise.

Evidence chain

Failed logins in the alert window

logins table

Period3 April 2024, 00:00 to 06:00

Failed attempts19

Accounts involved4

Context, not cited

YOUR QUERIES

1. Find the failed logins

Written SQL · 1 run

SELECT event_id, ts, account_id, source_ip, location FROM logins WHERE result = 'failure' AND ts >= '2024-04-03 00:00:00' AND ts < '2024-04-03 06:00:00' ORDER BY ts

Filtering: WHERE keeps the rows that meet every condition joined with AND. Text timestamps in YYYY-MM-DD HH:MM:SS form compare in time order.

2. Count attempts by account

Written SQL · 1 run

SELECT account_id, COUNT(*) AS attempts FROM logins WHERE result = 'failure' AND ts >= '2024-04-03 00:00:00' AND ts < '2024-04-03 06:00:00' GROUP BY account_id

MISSED INFORMATION

No unsupported claims on the way. Every finding followed from a result.

What these records do not show

Who was using the device. The logs record an account, not a person.

How the correct password was obtained.

Whether the downloaded files left the device.

Why alex.chen also signed in from the office at 09:10 on the same day. Ask the account owner.

Why I am showing you the gaps

Because closing them is the job I am asking to do. None of it is research. It is review, integration and content, and I know exactly what each game needs.

RÉMI COUTURE · GAMES FOR PROJECT ARES

22

From eight prototypes to games in the hub, with the team.

DAYS 1 TO 30

Get three in

- Agree the host contract with the platform team and write the hub adapter.
- SME review of Network Defense, Query Hunt and Cloud Command, using the sheets that exist.
- Map those three to NICE work roles and KSAs with the content team.
- Confirm or replace the borrowed art.

DAYS 31 TO 60

Test and widen

- Internal playtest. Measure completion, replays and common mistakes, then fix what confuses people.
- Review sheets and SME pass for Phishing, Triage and SOC.
- Japanese strings for the first three.
- A second case for the game learners replay most.

DAYS 61 TO 90

Finish and extend

- Dispatch and Privilege Path through review and into the hub.
- First new game from the design guide's Tier 1 list, built with the games team.
- An authoring guide so content developers can add scenarios themselves.
- A short report on what the playtests showed.

These are targets, and they depend on SME and platform time I do not control. I would rather agree the order with you than promise a date I cannot own.

This is working software, not a proposal. Starting again would mean rebuilding five things at once.

PIECE	WHAT EXISTS TODAY	WITHOUT IT
Game design	Eight different mechanics, each one built so it cannot be won without the knowledge	Finding a designer who knows both games and cyber training
Authored content	8 phishing messages with 32 cues, 6 alerts with 29 evidence items, 6 dispatch reports with 38 questions, 5 SQL challenges on 2 datasets, 3 access puzzles, 3 cloud checkpoints, and an explanation for every important mistake	Writer and SME time for every scenario from a blank page
Engineering	About 22,000 lines of TypeScript, eight static builds, one shared pattern	A front-end build from an empty folder
Balance and proof	209 passing tests, including the ones that prove the lazy strategies lose	Finding the exploits the hard way, with real learners
Review and integration	Four SME review packs ready to sign, and a draft host contract with an activity ID per game	Starting both conversations from nothing

Ready to show a customer

Any of the eight can be demoed in a browser today. No range to start, no sign-in, nothing to install. That is a new thing to put in front of a prospect this quarter.

The rare part is the combination

Game design, UX, front-end build, and knowledge of Ares from the inside. My 89 commits in the platform cover the world map, dashboard, debrief and the HUD screens, so these games are built to sit in the product, not beside it.

One owner, low cost

This does not need a new team or new infrastructure. It needs the person who built it, some SME hours, and a small integration step.

I want to stay and build this with you. Here are three ways to make it work.

PREFERRED

Full time, games and learner experience

Current hours and rate

I move toward the games team. The 90-day plan runs as written, and I keep supporting the player UI I already own: world map, dashboard, debrief and the HUD screens.

REDUCED

Half time, games only

About 20 hours a week

One focus: take the eight to production in priority order, then new games. The plan stretches to roughly twice the calendar time. UI work is by request.

MINIMUM

Per game, fixed scope

Agreed game by game

You pick a game. I deliver it reviewed, integrated and tested for an agreed scope. No standing commitment, and the work continues at whatever pace you choose.

What you keep

Eight working games and the person who knows every rule in them. A designer with twenty years of game experience who also builds. Someone already inside the product: the world map, dashboard, debrief and 16 HUD screens are my work.

What stops without an owner

These games are not in the platform repository yet. Without someone to take them through review and integration, they stay prototypes, and the open slots in the hub stay open.

My ask: thirty minutes to play two of them with you, and a decision on which option fits. The games are ready to demo to a customer now; the only open question is who takes them the rest of the way.